

LLaMA 70B Inference Performance (M2 Max)

